Predicting Craigslist NYC Apartment Rental Prices with Linear Regression

Ami Tian

Project Goal: Do an apartment's features help predict its rental price?

The small counter. Cameron Knowlton

https://www.insider.com/worst-apartment-new-york-city-tiktok-reactions-2021-2

Web Scraping and Data Cleaning

- Selenium/BeautifulSoup for web scraping
- Features:
  - # of bedrooms (BR)
  - # of bathrooms (BA)
  - Location
  - Square footage (SQFT)
  - Pets allowed
  - Application Fee
  - Broker's Fee
  - Laundry
  - Parking
- Cleaning: Dealing with outliers, incorrect/misplaced information, removing duplicate posts

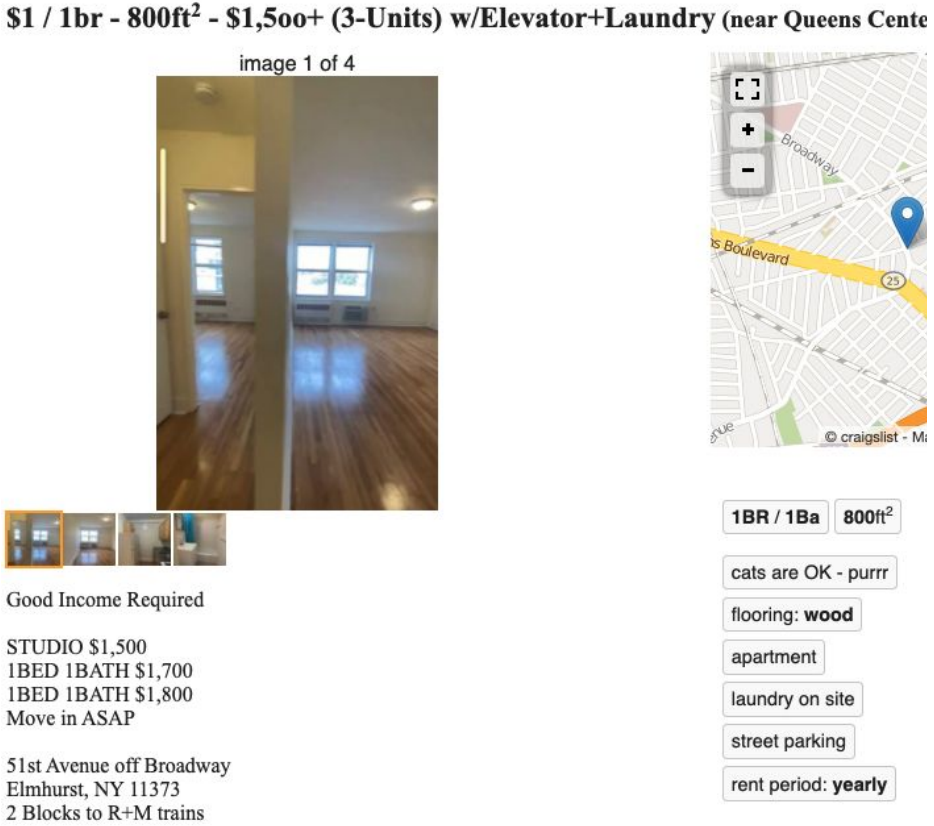

Data Prep and Feature Engineering

- Out of 4609 listings, only 969 listed square footage.
  - Imputed missing values based on median, borough, bedrooms.
- Created dummy variables for categoricals
  - Boolean: pets, application fee, broker's fee
  - Washer/Dryer: none, in-unit, in-building
  - Parking: none, off-street (includes carports and garages), valet
  - Borough: Brooklyn, Manhattan, Queens, the Bronx, Staten Island

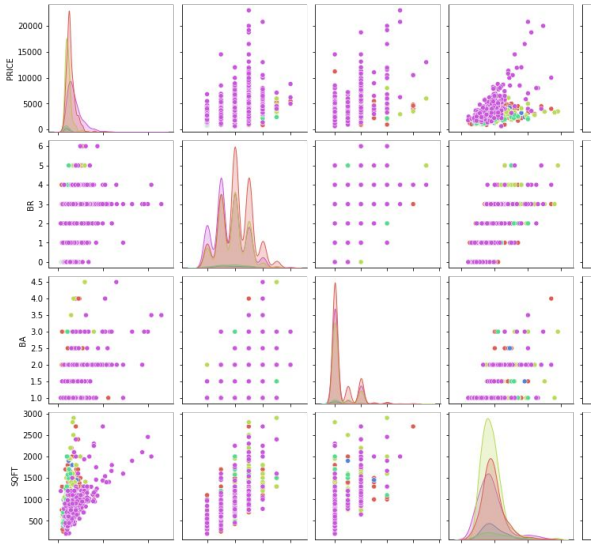

More Data Prep and Feature Engineering

- Removed some features (application fee, pets)
- Added interaction terms
  - BR / BA
  - SQFT * MANHATTAN

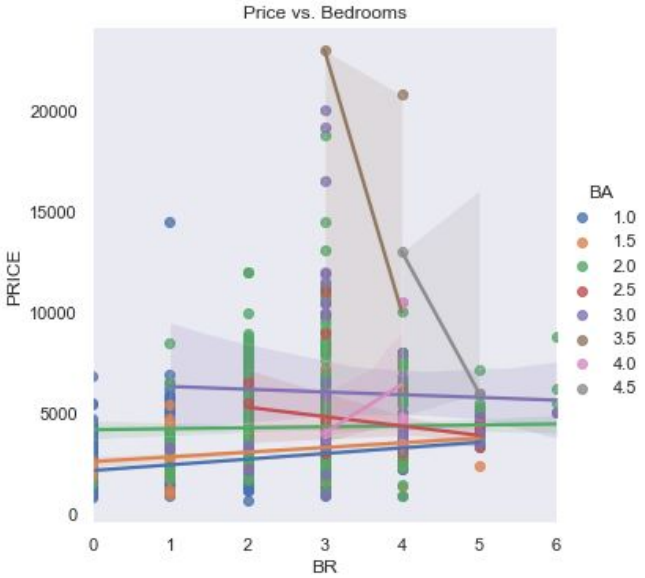

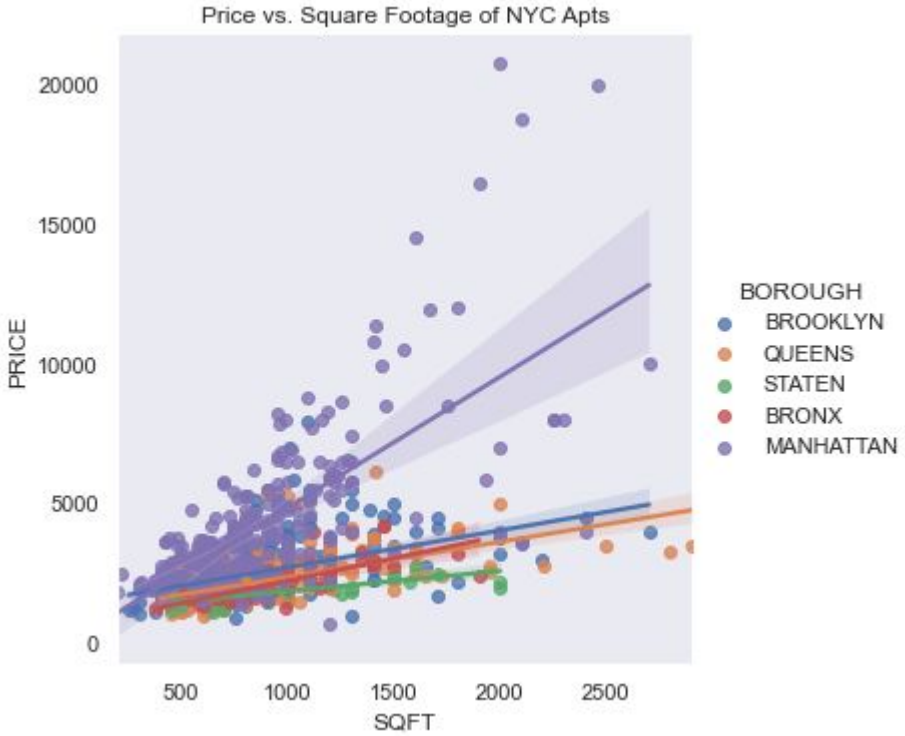

Issues

- Multicollinearity
- Heteroskedasticity

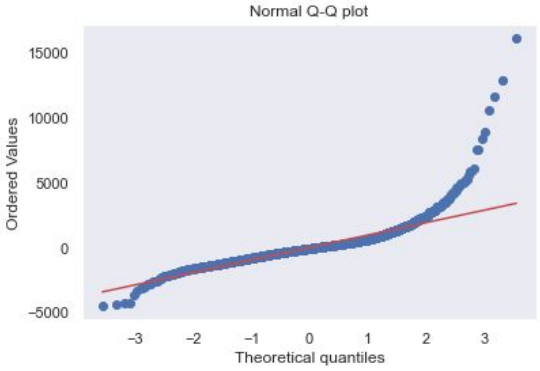

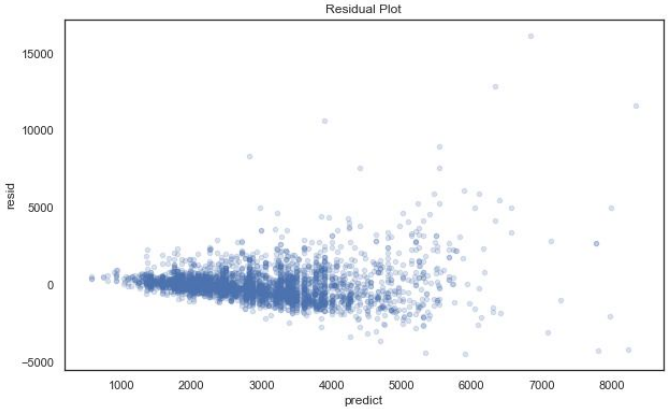

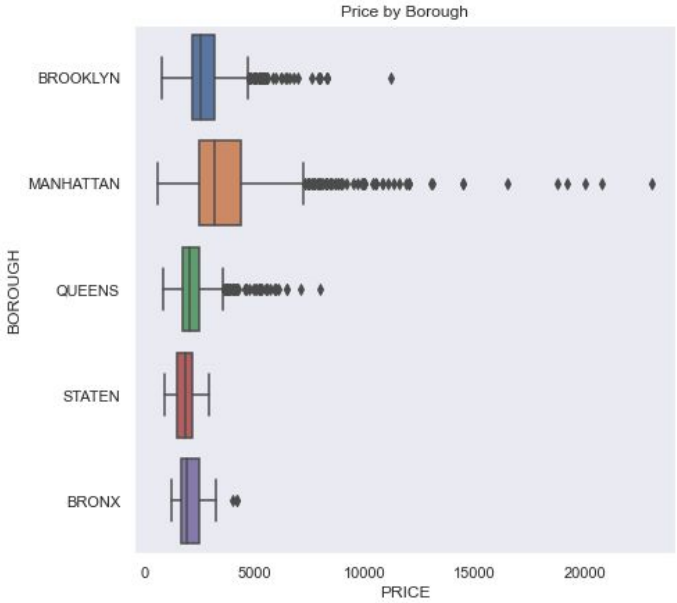

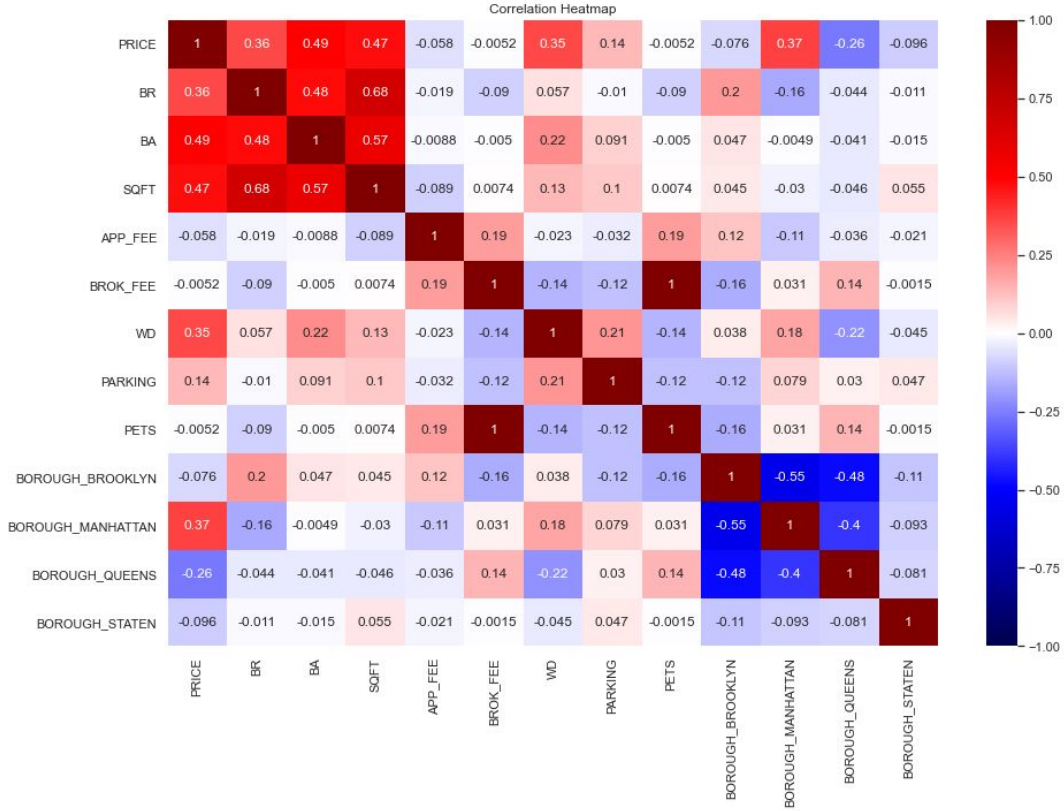

	variables	vif
0	BR	24.626210
1	BA	8.557337
2	SQFT	4.497295
3	BROK_FEE	1.078687
4	WD	1.184902
5	PARKING	1.087153
6	BOROUGH_BROOKLYN	19.939513
7	BOROUGH_MANHATTAN	29.935476
8	BOROUGH_QUEENS	16.101196
9	BOROUGH_STATEN	2.444164
10	MNH-SQFT	12.682606
11	BR-BA	17.592696

Model Selection and Results

- Used kfold cross-validation on training set to get the R^2 score, mean absolute error (MAE) and root mean squared error (RMSE)
- Applied Ridge and Lasso regularization, compared scores
- Lasso regularization improved scores slightly

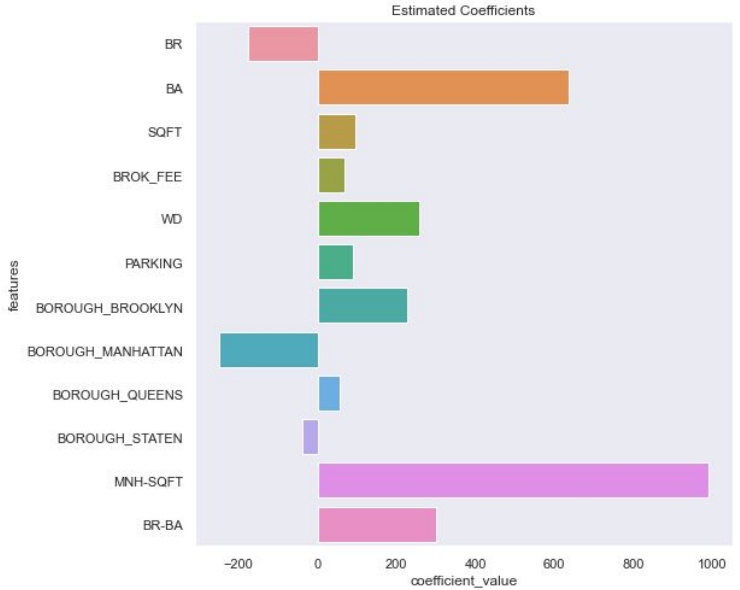

TRAINING DATA

- R^2 value: 0.526
- Mean Absolute Error: $648.65
- Root Mean Squared Error: $1,027.45

TEST DATA

- R^2 value: 0.519
- Mean Absolute Error: $671.15
- Root Mean Squared Error: $1,127.72

Conclusions

- The model can only explain about 50% of an apartment's rental price
- Other possible factors: neighborhood, amenities, utilities
- External factors: housing policies, market forces

Real Estate

New York Renters Face 70% Increases as Pandemic Discounts Expire

The era of widespread Covid concessions for apartment hu

NEW YORK

New York City, Income Inequality Capital Of America, Now Also Facing Soaring Rent Prices: Report

Future Steps

- Get more data, use information from other sites (e.g. StreetEasy, apartments.com)
- Look into more features
- Figure out more accurate interaction terms
- Categorize by neighborhood instead of borough, get more information about location
- Try log transformation or use weighted least squares regression to deal with heteroskedasticity